

# Using A2A with SWE-bench for Analyzing Agent Capability

Manish Muttreja

manishmuttreja@gmail.com

## Abstract

SWE-bench has emerged as the de facto standard for evaluating language model agents on real-world software engineering tasks. While recent systems like mini-SWE-agent achieve 65% resolution on SWE-bench Verified, the benchmark suffers from three critical limitations: (1) **data contamination**—models may have memorized repositories and patches during pretraining; (2) **patch-only scoring**—evaluation ignores the engineering process; and (3) **static test dependence**—fixed test suites can be overfit. I present **SWE-Bench-A2A**, an evaluation framework that quantifies these gaps through four key techniques: *retro-holdout mutations* [6] applied to SWE-bench to reveal contamination, *adversarial testing* (fuzz, edge case, mutation testing) that exposes patch fragility beyond test-pass metrics, *trajectory-based process scoring* capturing the full engineering workflow, and a *reproduction-first gate* enforcing understanding before patching. The experiments quantify technique impact: **retro-holdout mutations reveal 2.9% performance drop** on mutated instances (indicating memorization), while **adversarial testing shows patches achieve only 16–22% mutation robustness** despite passing repository tests. Cross-model evaluation with GPT-4.1, GPT-5.2, **Claude Sonnet 4.5** (27.7% avg), Claude Opus 4.1 (16.7%), and Claude 3 Haiku (18.9%) demonstrates framework generality across 5 models and 2 providers. These findings suggest current “resolved” metrics may overstate true capability by 3–5 $\times$ . I provide Dockerfiles and CI scaffolding for AgentBeats integration.

## 1 Introduction

The rapid advancement of large language models (LLMs) has enabled a new class of *software engineering agents*—systems that can understand codebases, diagnose bugs, and generate patches with minimal human intervention. Evaluating these agents requires benchmarks that capture the complexity of real-world software engi-

neering while resisting the pitfalls of static evaluation.

SWE-bench [1] represents a significant step forward, drawing from 2,294 real GitHub issues across 12 popular Python repositories. Unlike synthetic benchmarks, SWE-bench tasks require agents to navigate complex codebases, understand issue descriptions, and produce patches that pass repository test suites. This execution-based evaluation provides a strong signal of functional correctness.

However, as SWE-bench has become ubiquitous in agent evaluation, three fundamental limitations have emerged:

1. **Data Contamination:** The repositories in SWE-bench (Django, Flask, Scikit-learn, etc.) are among the most common in LLM training corpora. Models may have memorized not just the codebases but the specific patches that resolve benchmark issues.
2. **Patch-Only Scoring:** Current evaluation awards full credit for any patch that passes tests, ignoring whether the agent understood the problem. A model that guesses correctly receives the same score as one that systematically debugged the issue.
3. **Static Test Dependence:** Fixed test suites can be overfit through pattern matching without true understanding. Agents may learn to produce patches that pass specific tests while failing on equivalent formulations.

I present **SWE-Bench-A2A**, an evaluation framework that addresses these limitations through four key techniques:

- **Reproduction Gate:** Agents must first produce a failing test that reproduces the bug, demonstrating understanding before patching.
- **Process Scoring:** Beyond pass/fail, the framework captures full agent trajectories and computes multi-dimensional scores for correctness, process quality, efficiency, and adaptation.

- **Anti-Memorization:** I apply retro-holdout mutations [6] to SWE-bench, transforming codebases with semantic-preserving renames. A fresh issue harvester provides never-before-seen tasks.
- **Dynamic Testing:** Beyond repository tests, the framework supports fuzz testing, mutation testing, and adversarial probes to detect overfitting.

The framework implements the Agent-to-Agent (A2A) protocol, enabling modular composition of assessors (Green Agents) and participants (Purple Agents). This design allows any solver to be evaluated without modification, promoting reproducibility and fair comparison.

## 2 Related Work

### 2.1 Code Generation Benchmarks

Early code benchmarks like HumanEval [2] and MBPP [3] evaluate function-level generation from docstrings. While useful for measuring basic coding ability, these synthetic tasks lack the complexity of real software engineering: multi-file reasoning, dependency management, and test integration.

### 2.2 Repository-Level Evaluation

SWE-bench [1] pioneered repository-level evaluation using real GitHub issues. The SWE-bench Verified subset (500 instances) provides human-validated instances with clearer specifications. The official leaderboard<sup>1</sup> tracks state-of-the-art systems, with mini-SWE-agent achieving 65% resolved and SWE-agent 1.0 as the open-source SOTA. However, these “% resolved” metrics may not fully capture agent capability due to contamination and overfitting concerns. Concurrent work like DevBench [4] extends to multi-language settings.

### 2.3 Contamination and Memorization

Data contamination in LLM benchmarks has been extensively documented [5]. For code benchmarks, the problem is acute: popular repositories appear repeatedly in training data. Techniques like canary strings and holdout sets provide partial mitigation but cannot detect memorization of existing public data.

---

<sup>1</sup><https://www.swebench.com/>

### 2.4 Process-Aware Evaluation

Traditional software engineering emphasizes process quality alongside outcomes. Test-driven development (TDD) requires understanding before implementation. The reproduction gate operationalizes this principle for agent evaluation.

## 3 Limitations of Current SWE-bench

### 3.1 Data Contamination

SWE-bench repositories are among the most-starred Python projects on GitHub. Analysis suggests substantial overlap with common training corpora:

- Django: 76k+ stars, extensive documentation
- Flask: 66k+ stars, widely referenced in tutorials
- Scikit-learn: 58k+ stars, standard ML library

Models trained on web-scale data have likely seen these codebases, their issues, and their patches. Performance on “unseen” tasks may reflect recall rather than reasoning.

### 3.2 Patch-Only Evaluation

Current scoring treats all passing patches equally:

$$\text{Score} = \mathbb{1}[\text{all tests pass}] \quad (1)$$

This binary metric ignores:

- Whether the agent understood the bug
- The quality of the debugging process
- Efficiency of the solution path
- Ability to handle ambiguity

### 3.3 Static Test Overfitting

Repository test suites, while valuable, have fixed specifications. Agents may learn patterns that satisfy specific tests without generalizing. A patch that passes `test_user_login` may fail on semantically equivalent `test_account_authentication`.

## 4 SWE-Bench-A2A Design

### 4.1 A2A Protocol Architecture

The framework implements the Agent-to-Agent protocol with two actor types:

**Green Agent (Assessor)** Orchestrates evaluation: provisions environments, dispatches tasks, verifies solutions, computes scores.

**Purple Agent (Solver)** Attempts tasks: receives issue descriptions, explores codebases, generates patches.

Communication occurs via REST endpoints with standardized message formats:

```
# Task creation
POST /a2a/task
{
  "title": "Fix_bug_#1234",
  "description": "...",
  "resources": {"repo": "...", "commit": "..."}
}

# Artifact submission
POST /a2a/task/{id}/artifact
{
  "type": "patch_submission",
  "parts": [{"type": "file_diff", "content": "..."}]
}
```

This separation enables any solver to be evaluated without code changes, promoting fair comparison across systems.

### 4.2 Reproduction Gate

Before accepting patches, the framework requires agents to demonstrate bug understanding through reproduction:

---

#### Algorithm 1 Reproduction Gate Protocol

---

**Require:** Issue description  $I$ , environment  $E$

- 1: Agent submits reproduction script  $R$
  - 2: Execute  $R$  in unpatched  $E$
  - 3: **if**  $R$  does not fail **then**
  - 4:   **reject**: “Reproduction must fail before patch”
  - 5: **end if**
  - 6: Agent submits patch  $P$
  - 7: Apply  $P$  to  $E$
  - 8: Run full test suite
  - 9: **return** verification result
- 

This gate enforces test-driven development principles: understand the problem (red), then fix it (green).

### 4.3 Trajectory-Based Process Scoring

The framework captures complete agent trajectories and computes multi-dimensional scores:

$$\begin{aligned} S = & 0.35 s_{\text{correct}} + 0.20 s_{\text{process}} \\ & + 0.15 s_{\text{efficiency}} + 0.15 s_{\text{collab}} \\ & + 0.10 s_{\text{understand}} + 0.05 s_{\text{adapt}} \end{aligned} \quad (2)$$

where the scoring dimensions are:

Category	Weight	Description
Correctness	0.35	Tests pass, patch applies
Process	0.20	Systematic exploration
Efficiency	0.15	Token/time usage
Collaboration	0.15	Information requests
Understanding	0.10	Reproduction quality
Adaptation	0.05	Response to feedback

Table 1: Scoring dimensions and weights

### 4.4 Anti-Memorization Strategies

#### 4.4.1 Retro-Holdout Mutations

Following the retro-holdout methodology introduced by Haimes et al. [6], I apply semantic-preserving mutations to detect contamination:

- **Variable renaming:** `data` → `payload`
- **Function renaming:** `get_user` → `fetch_account`
- **Class renaming:** `UserManager` → `AccountHandler`
- **Comment perturbation:** Rephrase docstrings

Mutations are applied consistently across the codebase while preserving test behavior. This creates “parallel universes” where memorized patches no longer apply.

#### 4.4.2 Fresh Issue Harvesting

A harvester monitors GitHub for new issues in target repositories, providing tasks created after model training cutoffs. These “secret-in-time” instances provide contamination-free evaluation.

## 4.5 Dynamic Testing

Beyond repository tests, the framework provides hooks for (not enabled by default in reported runs):

**Fuzz Testing** Property-based tests with random inputs

**Mutation Testing** Assert patches handle code mutations

**Adversarial Probes** LLM-generated edge cases

## 5 Implementation

### 5.1 System Architecture

The implementation consists of several key components:

- **A2A Server:** FastAPI-based REST API implementing the A2A protocol with endpoints for task management, artifact submission, and health checks.
- **Environment Orchestrator:** Docker-based container management with JIT provisioning, repository cloning, and commit checkout.
- **Verification Engine:** Patch application, test execution with timeout handling, and flaky test detection.
- **Trajectory Capture:** Action logging with database persistence and streaming support.
- **LLM Solver:** Integration with OpenAI/Anthropic APIs for reproduction script and patch generation. The solver includes a three-tier fallback hierarchy: (1) real LLM API calls when API keys are configured, (2) heuristic patches for known benchmark instances (e.g., django-11099), and (3) mock responses when no API access is available. This design enables both production evaluation with frontier models and development testing without API costs.

### 5.2 Docker Images

I provide Dockerfiles for containerizing the Green and Purple agents. Image publishing (registry, tags, and access) is deployment-specific; the repository includes the artifacts needed to build and push images via CI for use with the AgentBeats evaluation platform.

## 6 Experiments

### 6.1 Setup

I ran six experiments to validate the framework and quantify the impact of each technique:

- **Experiments 1–4 (Baseline Performance):** Establish baseline metrics using semantic patch comparison across 3–100 instances with GPT-4o, GPT-4.1, and GPT-5.2 as Purple Agents.
- **Experiment 5 (Anti-Contamination—Key Result):** Retro-holdout mutation testing to quantify memorization vs understanding.
- **Experiment 6 (Adversarial—Key Result):** Fuzz, edge case, and mutation testing to quantify patch robustness beyond test-pass.

The first four experiments provide baseline metrics; experiments 5–6 quantify the impact of the novel techniques.

### 6.2 Experiment 1: Integration smoke test (3 Django instances)

Purpose: ensure end-to-end plumbing (environment provisioning, A2A dispatch, patch apply, test execution) works under Docker.

- django\_django-11099: UsernameValidator trailing newline (passed via heuristic baseline)
- django\_django-11133: HttpResponse charset handling (LLM patch failed to apply)
- django\_django-11179: model\_to\_dict for unsaved model (LLM patch failed to apply)

Instance	Patch	Tests	Time	Source
django-11099	✓	3/3	74s	Heuristic
django-11133	✗	0/0	73s	LLM
django-11179	✗	0/0	71s	LLM
<b>Total</b>	33.3%	-	-	-

Table 2: Integration smoke test: confirms Docker + A2A pipeline; highlights solver fragility on diff formatting.

**Important note:** The 33.3% success rate is from a heuristic baseline, not the LLM solver. The LLM-only success rate was 0% (0/2 tasks where LLM was used). This confirms infrastructure correctness while highlighting LLM solver fragility on diff formatting.

### 6.3 Experiment 2: GPT-4o benchmark (20 instances)

Purpose: measure solver quality with a stronger model on a broader slice. Evaluation uses **semantic patch**

**comparison** (code-change overlap) rather than strict line matching.

Metric	Value
Tasks Tested	20
Correct File Identification	100% (20/20)
Perfect Solutions (100% match)	25% (5/20)
High Match (>50%)	35% (7/20)
Average Semantic Match	43.2%
Composite Score $S$ (LLM-only)	0.43
Total Tokens	18,169
Total Cost	\$0.120
Cost per Task	\$0.006

Table 3: GPT-4o aggregate metrics on 20 SWE-bench Verified instances.

#### Representative outcomes (semantic match shown):

- **100%** `sklearn--sklearn-14141`: add `joblib` to `show_versions` deps (perfect semantic match).
- **100%** `django--django-13406`: `queryset` handling fix (perfect).
- **93%** `pallets--flask-5014`: blueprint registration fix (near-perfect).
- **80%** `sympy--sympy-23534`: symbol handling (strong partial).
- **0–50%** Several SymPy/Django tasks: correct file localization but partial or divergent semantics.

#### Key findings:

1. **File localization remains perfect:** 100% correct files on 20/20 tasks, confirming strong navigation.
2. **Semantic quality is mixed:** 25% perfect, 35% high-match; average semantic match rises to 43.2% on the larger slice.
3. **Repository difficulty:** Django and Flask skew higher (multiple 100%/93% cases); SymPy and some Django tests remain challenging with 0–50% matches.
4. **Cost efficiency holds:** \$0.006 per task with frontier model API calls.

**Context vs. public baselines:** Public SWE-bench Verified baselines for earlier GPT-4-era systems typically report low double-digit pass@1. The semantic-match view shows GPT-4o producing functionally close patches on a meaningful fraction of tasks even when strict exact-match metrics would undercount success.

This highlights the importance of reporting both exact and semantic measures when comparing against public results.

## 6.4 Experiment 3: Multi-model comparison (10 instances)

Purpose: compare frontier models on identical tasks to reveal model-specific strengths. Each model processed the same 10 SWE-bench Verified instances.

Metric	GPT-4o	GPT-4.1	GPT-5.2
Perfect (F1=100%)	<b>1</b>	0	0
High Match ( $\geq 50\%$ )	<b>2</b>	2	0
Files Correct	<b>10/10</b>	9/10	10/10
Avg F1 Score	<b>18.3%</b>	17.2%	7.0%
Cost	<b>\$0.063</b>	\$0.068	\$0.088

Table 4: Multi-model comparison on 10 identical SWE-bench tasks.

Instance	GPT-4o	GPT-4.1	GPT-5.2
sympy-22914	0%	<b>67%</b>	44%
sympy-23950	0%	<b>10%</b>	0%
sklearn-14141	<b>100%</b>	0%	0%
django-16082	0%	0%	0%
django-13406	<b>33%</b>	15%	7%
django-16429	0%	0%	0%
sympy-13757	0%	0%	0%
sympy-23534	0%	0%	0%
sympy-19040	0%	0%	0%
django-14534	50%	<b>80%</b>	18%

Table 5: Per-instance F1 scores across models (best per row in bold).

#### Key findings:

1. **GPT-4o leads at small scale:** Highest average F1 (18.3%) and only model with a perfect solution (sklearn-14141).
2. **Model-specific strengths:** GPT-4o uniquely solved sklearn-14141 perfectly; GPT-4.1 achieved highest score on sympy-22914 (67%) and django-14534 (80%).
3. **Consistent difficulty:** SymPy tasks (13757, 23534, 19040) remain challenging for all models (0% across the board).
4. **File localization robust:** 90–100% correct file identification across all models.

**Note on sampling variance:** GPT-5.2’s low score (7.0%) on this 10-task sample vs. its higher score (44.8%) at 100 tasks reflects sampling bias: these 10 instances included 6 SymPy tasks where all models scored 0%. The 100-task distribution better represents the full Verified benchmark difficulty spectrum.

## 6.5 Experiment 4: Large-scale benchmark (100 instances)

Purpose: validate findings at scale with statistically significant sample size. GPT-4.1 and GPT-5.2 each processed 100 SWE-bench Verified instances.

Metric	GPT-4.1	GPT-5.2
Tasks Completed	99/100	<b>100/100</b>
Files Correct	82.8%	<b>88.0%</b>
High Match ( $\geq 50\%$ )	36.4%	<b>40.0%</b>
Avg Semantic Match	38.7%	<b>44.8%</b>
Total Cost	\$1.30	<b>\$1.12</b>
Cost per Task	\$0.013	<b>\$0.011</b>

Table 6: 100-task benchmark: GPT-5.2 outperforms GPT-4.1 at scale.

### Key findings at scale:

- GPT-5.2 wins comprehensively:** Higher semantic match (44.8% vs 38.7%), more high-match solutions (40 vs 36), better file localization (88% vs 82.8%), and lower cost.
- Scale changes rankings:** At 10 tasks, GPT-4.1 led; at 100 tasks, GPT-5.2 dominates—demonstrating the importance of large-scale evaluation.
- Both models reliable:** 99–100% task completion shows production-ready robustness.
- Cost efficiency improves:** \$0.011–0.013 per task at scale vs \$0.006 in earlier runs reflects more complex tasks in the full distribution.

## 6.6 Experiment 5: Anti-Contamination Testing (Key Result)

**Purpose:** Quantify the contamination gap by comparing performance on original vs. retro-holdout mutated instances. Performance drops reveal memorization.

### Key findings:

- GPT-4.1 shows more contamination:** 2.9% performance drop on mutated instances vs GPT-5.2’s slight improvement (-0.8%).

Metric	GPT-4.1	GPT-5.2
Verified Avg Similarity	20.2%	21.6%
Mutated Avg Similarity	17.2%	22.3%
Performance Drop	2.9%	-0.8%
Avg Contamination Score	6.5%	<b>5.8%</b>
High Contamination (>30%)	7/100	7/100

Table 7: Anti-contamination at scale: GPT-4.1 shows higher contamination than GPT-5.2.

- Both models have similar high-contamination count:** 7/100 instances with >30% contamination.
- GPT-5.2 more robust to mutations:** Actually improved slightly on mutated instances, suggesting less reliance on memorization.
- Specific contamination:** `sklearn-14141` showed 100% contamination for GPT-4o (100%→0% on mutated version), demonstrating exact patch memorization.

## 6.7 Experiment 6: Adversarial Testing (Key Result)

**Purpose:** Quantify the robustness gap by testing patches against fuzz inputs, edge cases, and code mutations. This reveals fragility hidden by test-pass metrics.

Metric	GPT-4.1	GPT-5.2
Instances Tested	10	10
Pass Rate	40.0%	<b>60.0%</b>
Avg Fuzz Score	95.7%	<b>97.7%</b>
Avg Adversarial Score	40.0%	<b>44.0%</b>
Avg Mutation Score	16.0%	<b>22.0%</b>
Overall Adversarial Score	47.1%	<b>51.3%</b>

Table 8: Adversarial testing (10 instances): fuzz, edge case, and mutation robustness.

### Key findings:

- GPT-5.2 wins on adversarial robustness:** 60% pass rate vs 40% for GPT-4.1, with higher scores across all metrics.
- High fuzz resistance:** Both models show >95% fuzz test scores, indicating patches include defensive code patterns.
- Low mutation scores:** 16–22% mutation scores indicate patches may be fragile—tests would not catch many code mutations.

4. **Adversarial handling:** 40–44% adversarial scores suggest patches may not handle all edge cases (null inputs, boundary conditions).

The adversarial testing framework provides complementary signal to semantic match: a patch can be semantically correct but still fragile to edge cases or mutations.

## 6.8 Experiment 7: Claude Model Family Benchmark (100 instances)

**Purpose:** Evaluate Anthropic’s Claude model family on SWE-bench to compare against OpenAI models and validate framework generality across model providers.

Metric	Claude 3 Haiku	Opus 4.1	Sonnet 4.5
Tasks Completed	98/100	90/100	100/100
Avg Semantic Match	18.9%	16.7%	27.7%
High Match ( $\geq 70\%$ )	1	0	8
Perfect Match ( $\geq 95\%$ )	0	0	0
Total Tokens	92,810	114,760	103,511

Table 9: Claude model family comparison (100 instances each). Sonnet 4.5 leads.

## Notable high-match results for Claude Sonnet 4.5:

- **Best:** sklearn-12585 (91.3%), django-11163 (91.7%), xarray-4629 (85.5%)
- **Strong:** django-13670 (65.2%), pytest-7205 (66.9%), django-11451 (91.3%)

This cross-provider evaluation demonstrates the A2A framework’s generality—it can evaluate any LLM backend without modification.

## 6.9 Trajectory Analysis

For successful cases, the captured trajectory shows:

1. scenario_select -> instance_id
2. provision_environment -> [container_id]
3. dispatch_task -> [purple_task_id]
4. receive_artifact -> reproduction_script
5. receive_artifact -> patch_submission
6. verification -> passed (tests)

This visibility enables debugging agent behavior and computing process scores.

## 7 Impact of Novel Techniques

Model	Tasks	Avg Match	High ( $\geq 70\%$ )
GPT-5.2	100/100	44.8%	40
GPT-4.1	99/100	38.7%	36
Claude Sonnet 4.5	100/100	27.7%	8
Claude 3 Haiku	98/100	18.9%	1
Claude Opus 4.1	90/100	16.7%	0

Table 10: Cross-provider model rankings on SWE-bench (100 instances).

### Key findings:

1. **GPT-5.2 leads overall:** 44.8% avg match, significantly ahead of all Claude models.
2. **Claude Sonnet 4.5 leads Anthropic family:** 27.7% avg match with 8 high-quality solutions, outperforming both Opus 4.1 (16.7%) and Haiku (18.9%).
3. **Opus 4.1 underperforms:** Surprisingly, Opus 4.1 shows lower performance (16.7%) than the smaller Sonnet 4.5, possibly due to API errors (10/100 failed).
4. **Framework generality validated:** The A2A framework successfully evaluates 5 different models across 2 providers without modification.

The experiments reveal that standard SWE-bench metrics may significantly overstate true agent capability. I quantify the impact of each novel technique:

### Anthropic Contamination Gap: Retro-Holdout Impact

Comparing performance on original vs. mutated instances reveals memorization:

Metric	Original	Mutated	Gap
GPT-4.1 Avg Similarity	20.2%	17.2%	-2.9%
GPT-5.2 Avg Similarity	21.6%	22.3%	+0.8%
High Contamination ( $>30\%$ )	–	7/100	–

Table 11: Retro-holdout mutations reveal hidden contamination (100 instances per model).

**Interpretation:** GPT-4.1 shows a 2.9% performance drop on mutated instances, indicating memorization. GPT-5.2 is more robust (+0.8%). However, 7% of instances show  $>30\%$  drops for both models. Instance sklearn-14141 showed 100% contamination for GPT-4.0 (perfect → 0% on mutated version), demonstrating memorization of exact patches.

## 7.2 Robustness Gap: Adversarial Testing Impact

Standard test-pass metrics miss patch fragility:

Test Type	Avg Pass Rate	Gap from 100%
Fuzz Tests	96.7%	-3.3%
Adversarial Edge Cases	42.0%	<b>-58.0%</b>
Mutation Tests	19.0%	<b>-81.0%</b>

Table 12: Adversarial testing reveals hidden fragility (10 instances, avg of GPT-4.1 and GPT-5.2).

**Interpretation:** While patches show high fuzz resistance (96.7%), they fail on 58% of adversarial edge cases and 81% of code mutations. This suggests “resolved” patches may break under real-world usage variations.

## 7.3 Process Gap: Trajectory Analysis Impact

Binary pass/fail ignores engineering quality:

Scoring Approach	Avg Score	Variance
Binary (pass/fail)	25%	High
Semantic Match	44.8%	Medium
Process Score $S$	0.43	Low

Table 13: Multi-dimensional scoring captures nuance.

**Interpretation:** Process scoring differentiates agents that systematically debug (high  $s_{process}$ ) from those that guess correctly. This enables fairer comparison and identifies agents with transferable engineering skills.

## 7.4 Estimated True Capability

Combining these findings, I propose an adjusted capability estimate:

$$\text{Adjusted Capability} \approx \text{Resolved} \times (1 - \text{ContamGap}) \times \text{MutationRobustness} \quad (3)$$

For a hypothetical 65% resolved agent (using GPT-4.1’s 2.9% contamination gap and 19% mutation robustness):

$$\text{Adjusted} \approx 65\% \times 0.97 \times 0.19 \approx \mathbf{12.0\%}$$

**Caveat:** This is a rough estimate combining metrics from different experiment scales (100-task contamination, 10-task adversarial). The key insight is that current “% resolved” metrics may significantly overstate true robust capability—potentially by **3-5×**—when contamination and robustness are considered.

## 8 Evaluation Slices

I propose four evaluation slices for comprehensive assessment:

**Verified** Standard SWE-bench Verified instances

**Mutated** Retro-holdout transformed versions

**Fresh** Newly harvested issues ( $\leq 24\text{h}$  old)

**Adversarial** Instances with fuzz/mutation testing

Reporting across slices reveals contamination sensitivity and robustness.

## 9 Limitations and Future Work

### 9.1 Current Limitations

- **Python only:** Current implementation focuses on Python repositories
- **Model variance:** Performance varies significantly by both model and repository. At 100 tasks, GPT-5.2 leads overall but GPT-4.0 and GPT-4.1 excel on specific tasks. SymPy consistently challenges all models.
- **Semantic vs. exact matching:** The semantic comparison shows models often produce functionally equivalent patches that differ syntactically from expected solutions. Binary pass/fail evaluation may underestimate true capability.
- **Mutation coverage:** Retro-holdout not yet integrated in live evaluation flow
- **Dynamic test generation:** Fuzz/adversarial commands require per-repo configuration

### 9.2 Future Directions

1. Integrate additional frontier models (Claude 3.5 Sonnet, Gemini 2.0) for Purple agent comparison
2. Complete retro-holdout pipeline with semantic equivalence verification
3. Implement default fuzz command packs for common frameworks
4. Extend to multi-language evaluation (TypeScript, Rust)
5. Add visual/multimodal signals for UI-related bugs
6. Scale evaluation to full SWE-bench Verified (500+ instances)

## 10 Conclusion

SWE-Bench-A2A provides tools to quantify three critical gaps in current SWE-bench evaluation:

**1. Contamination Gap (Retro-Holdout):** The mutation testing reveals that 7% of instances show >30% performance drop when codebases are transformed with semantic-preserving renames. This indicates memorization of specific patches rather than true understanding. One instance showed 100% contamination—a “perfect” solution that completely failed on the mutated version.

**2. Robustness Gap (Adversarial Testing):** While patches pass repository tests, they fail on 81% of code mutations and 58% of adversarial edge cases. Current “% resolved” metrics may overstate true capability by 3–5× when robustness is considered.

**3. Process Gap (Trajectory Scoring):** Binary pass/fail treats lucky guesses equally with systematic debugging. The multi-dimensional scoring (correctness, process, efficiency, adaptation) provides fairer comparison that rewards transferable engineering skills.

These findings have implications for the SWE-bench leaderboard. While systems like mini-SWE-agent achieve 65% resolved, this analysis suggests **true contamination-adjusted, robustness-weighted capability may be significantly lower**. I recommend the community adopt:

- **Multi-slice reporting:** Report on Verified, Mutated, and Adversarial slices
- **Robustness metrics:** Include mutation score alongside pass rate
- **Contamination disclosure:** Flag instances with high verified-to-mutated drops

I release ready-to-run Docker images compatible with AgentBeats to encourage reproducible, process-aware benchmarking that rewards true engineering ability over memorization.

## Acknowledgments

I thank the SWE-bench team for their foundational work and the AgentBeats community for evaluation infrastructure.

## References

- [1] C. E. Jimenez, J. Yang, A. Wettig, S. Yao, K. Pei, O. Press, and K. Narasimhan. SWE-bench: Can language models resolve real-world GitHub issues? In *International Conference on Learning Representations (ICLR)*, 2024.
- [2] M. Chen, J. Tworek, H. Jun, Q. Yuan, H. Pinto, J. Kaptan, et al. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021.
- [3] J. Austin, A. Odena, M. Nye, M. Bosma, H. Michalewski, D. Dohan, et al. Program synthesis with large language models. *arXiv preprint arXiv:2108.07732*, 2021.
- [4] DevBench Team. DevBench: A comprehensive benchmark for software development. *arXiv preprint*, 2024.
- [5] O. Sainz, J. Campos, I. García-Ferrero, J. Etxaniz, O. Lopez de Lacalle, and E. Agirre. NLP evaluation in trouble: On the need to measure LLM data contamination for each benchmark. In *Findings of EMNLP*, 2023.
- [6] J. Haimes, B. Peng, and S. Mukherjee. Benchmark Inflation: Revealing LLM Performance Gaps Using Retro-Holdouts. *arXiv preprint arXiv:2402.14857*, 2024.
- [7] AgentBeats Platform. Agent registry for AI evaluation. <https://agentbeats.dev/>, 2024.

## A A2A Protocol Specification

### A.1 Agent Card Format

```
1 {  
2     "name": "SWE-bench Green Agent",  
3     "version": "1.0.0",  
4     "agent_id": "uuid",  
5     "capabilities": ["swebench_evaluation"],  
6     "endpoints": {  
7         "task": "/a2a/task",  
8         "health": "/health"  
9     }  
10 }
```

### A.2 Artifact Types

- reproduction\_script: CODE artifact with failing test
- patch\_submission: FILE\_DIFF artifact with unified diff

- `assessment_result`: JSON artifact with verification results

## B Scoring Formula Details

### B.1 Correctness Score

$$s_{\text{correct}} = 0.6 \cdot \mathbb{I}[\text{pass}] + 0.3 \cdot \frac{\text{tests\_passed}}{\text{total\_tests}} + 0.1 \cdot \mathbb{I}[\text{patch\_applied}] \quad (4)$$

### B.2 Process Score

$$s_{\text{process}} = 0.4 \cdot s_{\text{exploration}} + 0.3 \cdot s_{\text{reasoning}} + 0.3 \cdot s_{\text{reproduction}} \quad (5)$$

### B.3 Efficiency Score

$$s_{\text{efficiency}} = 0.4 \cdot \frac{T_{\text{budget}} - T_{\text{used}}}{T_{\text{budget}}} + 0.4 \cdot \frac{N_{\text{budget}} - N_{\text{tokens}}}{N_{\text{budget}}} + 0.2 \cdot \frac{1}{1 + \text{attempts}} \quad (6)$$